

# Single Sample and Paired Samples $t$ tests

Dr. J. Kyle Roberts

Southern Methodist University  
Simmons School of Education and Human Development  
Department of Teaching and Learning

## Confidence Intervals for the Mean

- Suppose that we have a sample of IQ scores from 10 students.
- How might we know if the mean IQ of our 10 students is statistically significantly different from the population mean IQ of 100?

```
> iqdata <- c(100, 110, 95, 98, 120, 118, 143, 94,  
+           82, 114)  
> mean(iqdata)
```


```
[1] 107.4
```

```
> sd(iqdata)
```

```
[1] 17.35383
```


## IQ Data

We can then plot this data with a horizontal line representing the population IQ score at 100.



## Sampling from IQ Data

Suppose that we were to draw 25 random samples from the population of all IQ scores and draw 95% confidence intervals around each mean. It would stand to reason that 95% of the means would have confidence intervals that captured the population mean. For example:



## More Study 1/Study 2 for Confidence Intervals

Suppose that we have two different datasets with the following properties and a population mean of 20:

	Study 1	Study 2
mean	10	15
SD	10	10
n	30	30

### Thought Question

Given the properties of the above two studies, which study would be MORE likely to have confidence intervals which did not capture the population mean?

## More Study 1/Study 2 for Confidence Intervals

Suppose that we have two different datasets with the following properties and a population mean of 20:

	Study 1	Study 2
mean	10	10
SD	5	10
n	30	30

### Thought Question

Given the properties of the above two studies, which study would be MORE likely to have confidence intervals which did not capture the population mean?

## More Study 1/Study 2 for Confidence Intervals

Suppose that we have two different datasets with the following properties and a population mean of 20:

	Study 1	Study 2
mean	10	10
SD	10	10
n	300	30

### Thought Question

Given the properties of the above two studies, which study would be MORE likely to have confidence intervals which did not capture the population mean?

## On Your Own

Try a dataset on your own. Consider the following data:

- $X = 7, 8, 7, 6, 4, 2, 3, 8$

Complete the following tasks:

- Calculate confidence intervals (by hand).
- Determine if our data is statistically significantly different from a population mean of 10.
- Compute the confidence intervals in R by `ci(x)`.



## From CIs to $t$ -calcs

- Recall from the Confidence Interval slides that we are looking to see if our sample mean captures the population mean given certain parameters.
- Rather than looking to see if our sample mean captures the population mean, we can compute the probability that our sample mean is different from the population mean.

## One-sample $t$ tests

- Occasionally we have a single sample of measurements and wish to evaluate a claim about the population mean,  $\mu$ , relative to some fixed, known value,  $\mu_0$ .
- If we assume that the sample,  $x_1, x_2, \dots, x_n$  is a random sample (i.e. the values of independent and identically distributed random variables) from a distribution with mean  $\mu$  and variance  $\sigma^2$  then, by the Central Limit Theorem, the random variable  $\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$  tends to a  $\mathcal{N}(0, 1)$  distribution for large  $n$ .
- The quantity  $\sigma/\sqrt{n}$  is called the standard error of the mean (SEM). It describes the amount of variability we expect to see in  $\bar{X}$ .
- Unfortunately, we do not know  $\sigma$ .
- The empirical SEM is  $s/\sqrt{n}$ . That is, it uses the estimate  $s$  of the (unknown) standard deviation  $\sigma$ .

## One-sample $t$ tests (continued)

- If the data,  $x_1, x_2, \dots, x_n$  are a random sample from a *normal* distribution,  $\mathcal{N}(\mu, \sigma^2)$ , then  $\frac{\bar{X} - \mu}{s/\sqrt{n}}$  has a  $T$  distribution with  $n - 1$  *degrees of freedom*.
- The  $T$  distribution is similar to the standard normal distribution but a bit more “spread out” or over-dispersed relative to the standard normal.
- In general we write the degrees of freedom as  $df$ . For a one-sample  $t$ -test  $df = n - 1$  but for other tests we will have different formulas for  $df$ .
- The over-dispersion of the  $T$  relative to the standard normal is greater for smaller values of  $df$ .
- For very large values of  $df$ , the  $T$  is essentially the same as the standard normal.

## $p$ -value of a statistical hypothesis test

- The  $p$ -value of a statistical hypothesis test is the probability of seeing the data that we did, or something even more unusual, if  $H_0$  were true.
- However,  $p$ -values are somewhat akin to pregnancy tests. You either are or you are not.
- Some would argue that a very small  $p$ -value is strong evidence against  $H_0$ .
- In some disciplines predetermined “significance levels”, such as 5% and 1%, are used. Thus a  $p$ -value less than 0.05 is “statistically significant”.
- For the single-sample  $t$  test, the  $H_0 : \mu = \overline{X}$ .

## One-sided versus two-sided alternatives

- Depending on the context of the data, we may be interested in an increase only ( $H_a : \mu > \mu_0$ ), or a decrease only ( $H_a : \mu < \mu_0$ ) or either type of a shift ( $H_a : \mu \neq \mu_0$ ).
- An increase only or a decrease only are called one-sided alternatives.
- Either type of shift ( $H_a : \mu \neq \mu_0$ ) is called a two-sided alternative.
- The intention to perform a one-sided test should be clear before the study is conducted and the data gathered. It should not be decided after examining the data.
- In many cases shifts in only one direction are desirable (we only wish to increase the yield, not decrease it; we only wish to decrease the proportion of defectives, not increase it) and the only sensible test is one-sided.

## Assumptions on the $t$ -test

- For the results of `t.test` to be valid, we must assume that we have a random sample (i.e. independent observations with an identical distribution) **and** that the underlying distribution is a normal distribution.
- If the sample size,  $n$ , is large we can relax the requirement of a normal distribution to “sort-of like a normal”.
- For small samples departures from normality become more important. However, when  $n$  is small we cannot check reliably for departures from normality.

## Practical Example for IQ Data

Remember again our IQ Data

```
> iqdata <- c(100, 110, 95, 98, 120, 118, 143, 94,  
+           82, 114)  
> t.test(iqdata, mu = 100)
```

One Sample  $t$ -test

data: iqdata

$t = 1.3485$ ,  $df = 9$ ,  $p\text{-value} = 0.2105$

alternative hypothesis: true mean is not equal to 100

95 percent confidence interval:

94.98581 119.81419

sample estimates:

mean of  $x$

107.4

- This means that our sample mean of 107.4 is not statistically significantly different from the population mean of 100 at the  $\alpha = 0.05$  level on 9  $df$ .

## Another IQ Data Example

This time consider the following dataset

```
> iqdata2 <- c(100, 110, 145, 115, 120, 118, 143,  
+           94, 122, 114)  
> t.test(iqdata2, mu = 100)
```

One Sample  $t$ -test

data: iqdata2

$t = 3.5325$ ,  $df = 9$ ,  $p\text{-value} = 0.006389$

alternative hypothesis: true mean is not equal to 100

95 percent confidence interval:

106.5089 129.6911

sample estimates:

mean of  $x$

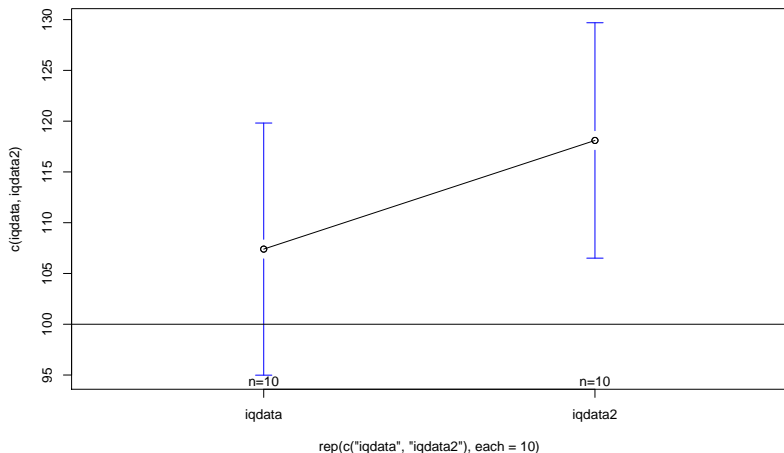
118.1

- This means that our sample mean of 118.1 *is* statistically significantly different from the population mean of 100 at the  $\alpha = 0.05$  level on 9  $df$ .



## Graphical Examples of `iqdata` and `iqdata2`

```
> plotmeans(c(iqdata, iqdata2) ~ rep(c("iqdata",  
+      "iqdata2"), each = 10))  
> abline(h = 100)
```



## Single Sample and Paired Samples $t$ test

- In the single sample  $t$  test we are testing the null:

$$H_0 : \mu = \overline{X}$$

- It may be that we have a different data scenario where we have a set of scores gathered at two different time points.
- This is very common in education studies where some assessment instrument is given as a pre-test and post-test with an intervention between the two data gatherings.
- With a paired samples  $t$  test, we can compute the statistical probability that the mean score at one time point is statistically significantly different from the mean score at another time point. Or:

$$H_0 : \overline{X}_{pre} = \overline{X}_{post}$$

## Paired Samples $t$ test Generalization

- Our original null hypothesis for the paired samples  $t$  test is just a generalization of the single sample  $t$  test null in that

$$\begin{aligned}H_0 : \bar{X}_{pre} &= \bar{X}_{post} \\ \bar{X}_{pre} - \bar{X}_{post} &= 0 \\ \bar{X}_{diff} &= 0\end{aligned}$$


- This last form of the null hypothesis now takes on the same form as a single sample  $t$  test where we are assuming the population mean difference score is 0.

## A Practical Example

```
> pre <- c(60, 70, 65, 68, 72, 49, 59, 82)
> post <- c(71, 79, 82, 78, 85, 60, 63, 97)
> c(mean(pre), sd(pre), mean(post), sd(post))
```

```
[1] 65.625000  9.898593 76.875000 12.040853
```

```
> plotmeans(c(pre, post) ~ rep(c("pre", "post"),
+   each = 8))
```





rep(c("pre", "post"), each = 8)

## Assumptions of the Paired Samples $t$ test

- Normality in the population.
- No interaction between pairs and treatment. This can happen when there is a ceiling effect in your data.
- Outliers



## Assumptions of the Paired Samples $t$ test cont.

```
> print(p1 <- xyplot(post ~ pre), split = c(1, 1,  
+      2, 1), more = TRUE)  
> print(tmd(p1), split = c(2, 1, 2, 1))
```





# Assumptions with Interaction Data

```
> pre1 <- seq(10, 100, by = 10)
> post1 <- seq(82, 100, length = 10)
> print(p2 <- xyplot(post1 ~ pre1), split = c(1,
+      1, 2, 1), more = TRUE)
> print(tmd(p2), split = c(2, 1, 2, 1))
```



# Assumptions with Outlier Data

```
> pre2 <- 1:10
> post2 <- c(11:19, 50)
> print(p3 <- xyplot(post2 ~ pre2), split = c(1,
+       1, 2, 1), more = TRUE)
> print(tmd(p3), split = c(2, 1, 2, 1))
```






## Practical Example continued

```
> cbind(pre, post, diff = pre - post)
```

	pre	post	diff
[1,]	60	71	-11
[2,]	70	79	-9
[3,]	65	82	-17
[4,]	68	78	-10
[5,]	72	85	-13
[6,]	49	60	-11
[7,]	59	63	-4
[8,]	82	97	-15

```
> plotmeans(pre - post ~ rep("difference", 8))
```



## $t$ .test in $R$

```
> t.test(pre, post, paired = T)
```

Paired t-test

data: pre and post

$t = -8.0452$ ,  $df = 7$ ,  $p\text{-value} = 8.792e-05$

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-14.55655 -7.94345

sample estimates:

mean of the differences

-11.25

```
> t.test(pre - post, mu = 0)
```

One Sample t-test

data: pre - post

$t = -8.0452$ ,  $df = 7$ ,  $p\text{-value} = 8.792e-05$

alternative hypothesis: true mean is not equal to 0

95 percent confidence interval:

-14.55655 -7.94345

sample estimates:

mean of x

-11.25